

CSE 486/586 Distributed Systems

Distributed Shared Memory

Steve Ko
Computer Sciences and Engineering
University at Buffalo

Overview

- Today: distributed shared memory, starting from some background on memory sharing
- Memory sharing **for a single machine**
 - Threads and processes
- Memory sharing **for different machines**
 - Threads and processes

Why Shared Memory?

- For sharing data
- There are two strategies for data sharing.
 - Message passing
 - Shared memory
- Message passing
 - Send/receive primitives
 - Explicit sharing → no synchronization (locks) necessary
- Shared memory
 - Memory read/write primitives (in your code, you could use regular variables)
 - Typically requires explicit synchronization (locks)
- Which is better?
 - Depends on your use case.
 - Multiple writers: perhaps message passing
 - (Mostly) read-only data: shared memory

Memory Sharing for Threads

- Threads belong to a single process, so **all threads share the same memory address space**.
- E.g., Java threads


```
class MyThread extends Thread {  
    HashMap hm;  
    MyThread(HashMap _hm ) {  
        this.hm = _hm;  
    }  
    public void run() {  
        ...  
        hm.put(key, value);  
    }  
}
```

```
HashMap hashMap = new HashMap();  
MyThread mt0 = new MyThread(hashMap); // hashMap is shared  
MyThread mt1 = new MyThread(hashMap);  
mt0.start();  
mt1.start();
```

Memory: Threads vs. Processes

- For threads, there's no special mechanism necessary to share memory.
- But, a process has **its own address space**, so by default, different processes do not share memory.
- Processes (on the same machine) can share memory regions with support from their OS.

Shared Memory on a Single Machine

- Shared memory is part of IPC (Inter-Process Communication).
 -  – What are other IPC mechanisms?
 - Files, (domain) sockets, pipes, etc.
- Shared memory API (POSIX C)
 - `shm_open()`: create and open a new object, or open an existing object. The call returns a file descriptor.
 - `mmap()`: map the shared memory object into the virtual address space of the calling process.
 - ...and others
- Semaphore API (POSIX C)
 - `sem_open()`: initialize and open a named semaphore
 - `sem_wait()`: lock a semaphore
 - `sem_post()`: unlock a semaphore
 - ...and others

Shared Memory Example* (in C)

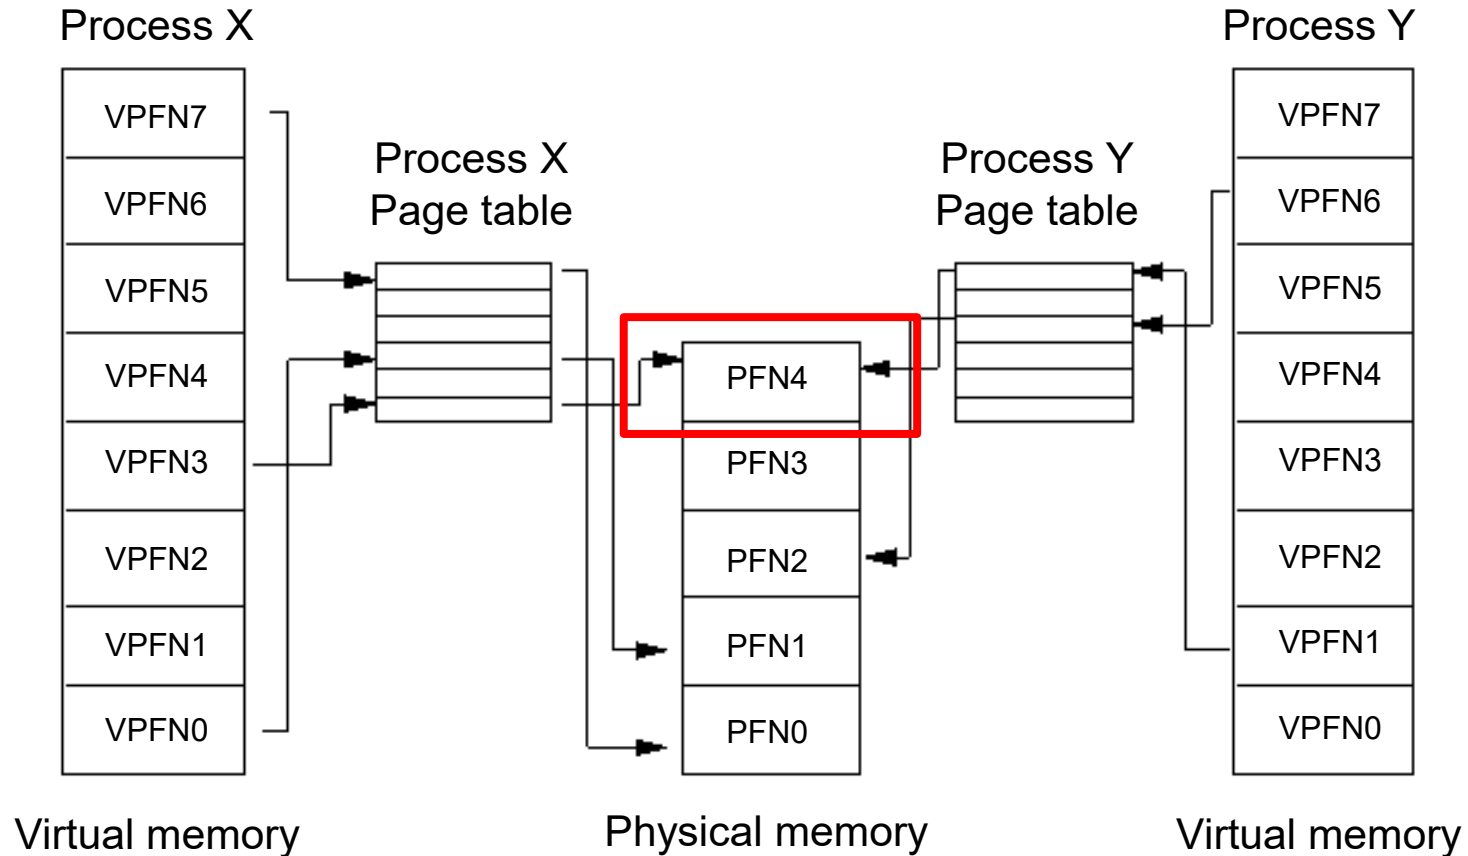
```
int main() {
    const char *name = "shared"; // shared with other processes
    int shm_fd;
    void *ptr;

    /* create the shared memory segment. name is shared. */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    ...
    /* now map the shared memory segment in the address space of
       the process */
    ptr = mmap(0, SIZE, PROT_READ | PROT_WRITE,
               MAP_SHARED, shm_fd, 0);

    sprintf(ptr, "%s", message0);

    return 0;
}
```

Shared Memory Implementation

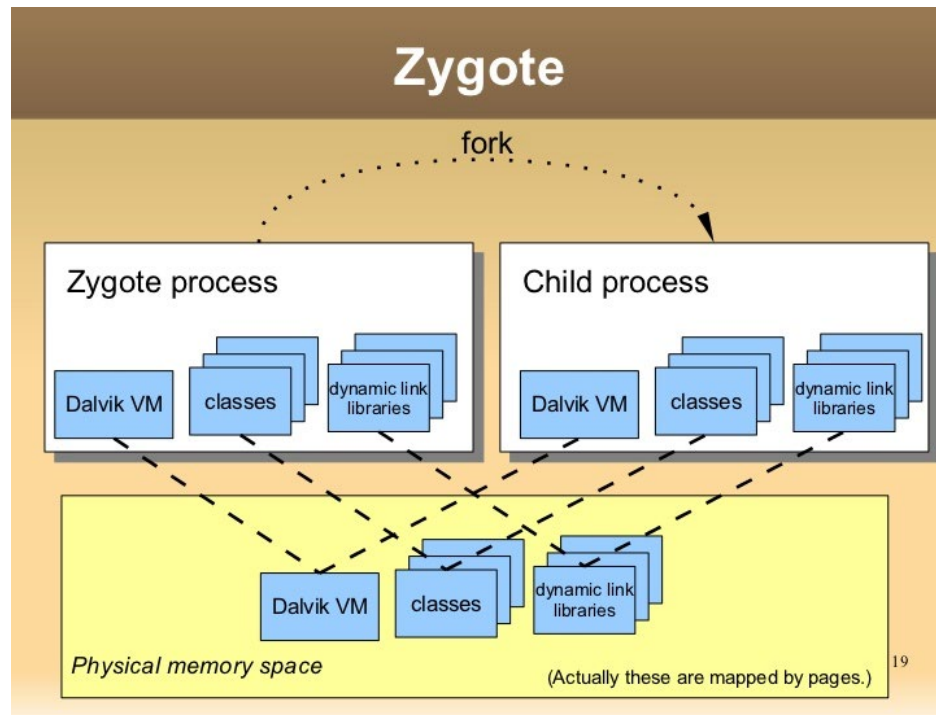


- VPFN: Virtual page frame number
- PFN: Physical page frame number
- Adapted from <http://tldp.org/LDP/tlk/mm/memory.html>

Shared Memory Use Case: Android

- All apps need framework API libraries, Java VM, etc.
 - Too expensive if all app processes have them in their memory space individually.
- Zygote: A process that starts everything else.
 - All app processes share memory with Zygote.

Image source: https://www.slideshare.net/tetsu.koba/android-is-not-just-java-on-linux/19-Zygote_forkZygote_process_Child_process



CSE 486/586 Administritivia

- PA3 grades will be posted today.
- PA4 deadline: 5/10
 - Please start early. The grader takes a long, long time.
- Survey & course evaluation
 - Survey: <https://forms.gle/eg1wHN2G8S6GVz3e9>
 - Course evaluation:
<https://www.smartevals.com/login.aspx?s=buffalo>
- If **both** have 80% or more participation,
 - For each of you, I'll take the better one between the midterm and the final, and give the 30% weight for the better one and the 20% weight for the other one.
 - (Currently, it's 20% for the midterm and 30% for the final.)
- No recitation today; replaced with office hours

Distributed Shared Memory


- We will discuss two cases.
 - DSM for processes
 - DSM for threads
- DSM for processes: different processes running on different machines sharing a memory page.
- The shared memory page is **replicated and synchronized** across different machines.
 - However, replication is not the goal (e.g., we're not keeping replicas to deal with failures).
- A generic way of doing this is at **the OS layer**.
 - Similar to the diagram on slide #8, but with processes on different machines

DSM Synchronization Options

- Write-update
 - A process updates a memory page.
 - The update is **multicast** to other replicas.
 - The multicast protocol determines consistency guarantees (e.g., FIFO-total for sequential consistency).
 - **Reads are cheap** (always local), but **writes are costly** (always multicast).
- Write-invalidate
 - Two states for a shared page: **read-only** or **read & write**
 - » Read-only: the memory page is **potentially replicated** on two or more processes/machines
 - » Read & write: the memory page is **exclusive** for the process (no other replica)
 - If a process intends to write to a read-only page, **an invalidate request is multicast** to other processes.
 - Later writes can take place **without communication (cheap)**.
 - Writes are only propagated **when there's a read by another process (cheap for write, costly for read)**.
 - But a write can be **delayed by invalidation (costly for write)**.

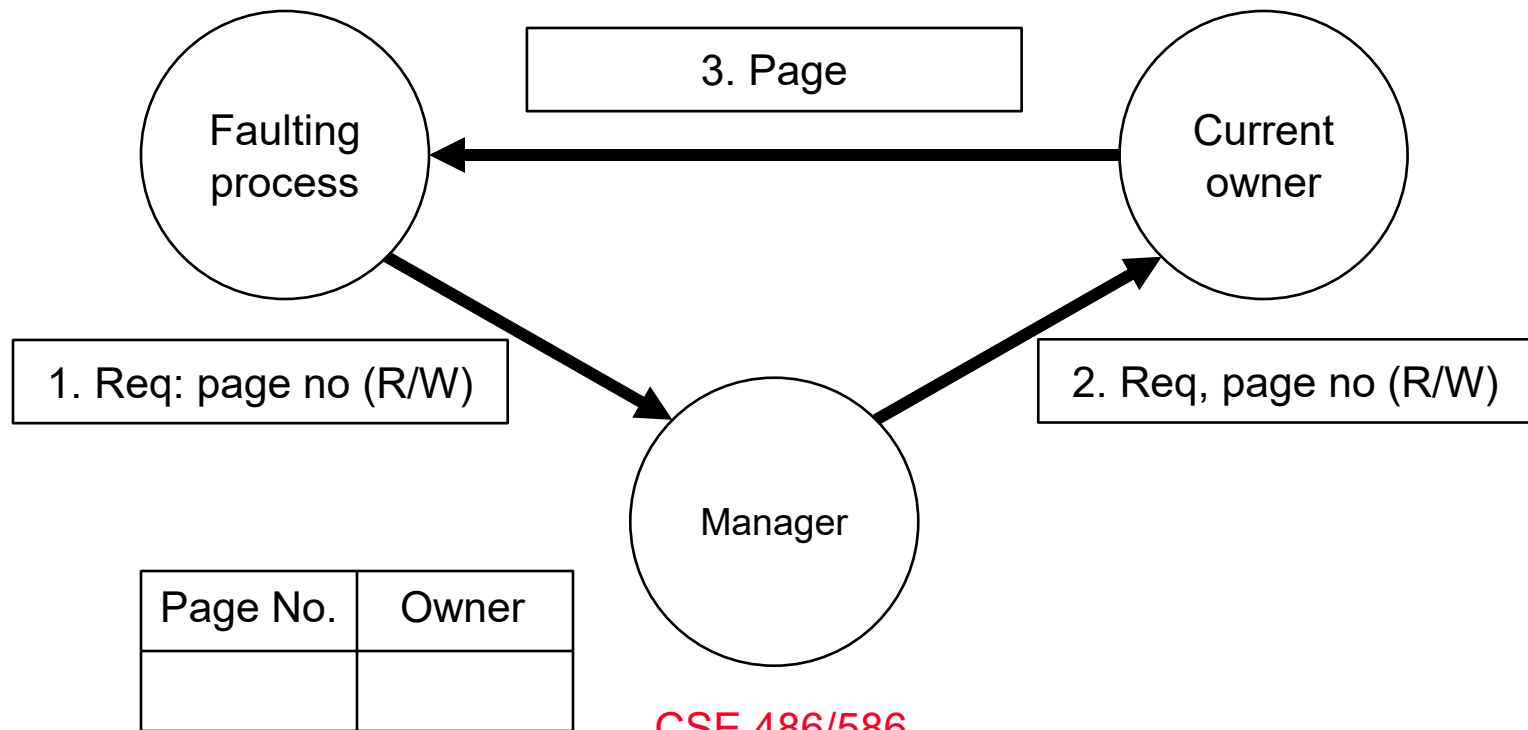
Write Invalidate Protocol Example

- Note: R fault and W fault can occur at any process



Example System: Ivy

- Implements a write-invalidation protocol
 - **Owner** of a page: the process with the most up-to-date
 - **Copyset** of a page: the processes with a replica
 - A **centralized manager** maintains ownership info.




Granularity Problem

- Let's assume that we operate at the page-level.
 - (But other implementations also have similar problems.)
 - Just as a reference, a Linux memory page is 4KB.
- Problem
 - When two processes (on two different machines) share a page, **it doesn't always mean that they share everything on the page.**
 - E.g., one process reads from and writes to a variable X, while the other process reads from and writes to another variable Y. If they are in the same memory page, the processes are sharing the page.

Granularity Problem

- True sharing
 - Two processes share **the exact same data**.
- False sharing
 - Two processes **do not share the exact same data**, but they access **different data from the same page**.



- False sharing problems
 - Write-invalidate: unnecessary invalidations
 - Write-update: unnecessary data transfers

Granularity Problem

- Bigger page sizes
 - Better handling for **updates of large amounts of data** (good)
 - **Less management overhead** due to a smaller number of units/pages to handle (good)
 - More possibility for **false sharing** (bad)
- Smaller page sizes
 - The opposite of the above
 - If there is an update of a large amount of data, it'll be broken down to many small updates, which leads to **more network overhead** (bad)
 - A smaller page size means more pages, which leads to more management overhead, i.e., **more tracking of reads and writes** (bad)
 - **Less** possibility of false sharing (good)

Thrashing


- Thrashing could happen with write-invalidate protocols.
- *Thrashing is said to occur when DSM spends **an inordinate amount of time invalidating and transferring shared data** compared with the time spent by application processes doing useful work.*
- This occurs when **several processes compete for a data item** or **for falsely shared data items**.

Thrashing

- Common scenario: **producer-consumer pattern**
 - Data is produced by a process and used by another process.
 - The producer will keep invalidating the consumer & the consumer will keep transferring data from the producer.
 - Write-update is better for this pattern.
- Solutions to thrashing
 - Manual avoidance: a programmer avoids thrashing patterns.
 - Timeslicing: once a process gains a write access to a page, **it retains it for a period of time**. Other processes' read/write requests are **buffered** during that period.


DSM for Threads

- Memory sharing among threads on different machines.
- Use case: **code (thread) offloading** from a smartphone to a server
 - Low-power smartphones augmented by high-power servers (computation & energy)
 - In some sense, it's done already (cloud backend), but DSM allows it without any programmer effort.



Example: Comet*

- Comet allows thread offloading for Android apps in Java
- Comet synchronizes the entire Java VM state.



*<https://www.usenix.org/conference/osdi12/technical-sessions/presentation/gordon>

Java Code Execution Background

- Memory: program code, stack, heap, & CPU state
- Stack & heap
 - Generally, the program stack handles **statically allocated objects** & method call **return addresses**.
 - The heap is used for **dynamically** allocated objects.

```
public class Ex {  
    public void method() {  
        int i = 0; // stack  
        HashMap hm = new HashMap(); // heap  
    }  
}
```

- CPU state
 - Android Java VM uses **registers** for instruction execution.
 - **The program counter (PC)** points to the next instruction to execute.
- For program execution, Java VM has an **execution loop**.
 - Fetches the next instruction that the PC points to.
 - Executes the new instruction
 - While executing, it uses registers, the stack, and the heap.

Comet Thread Migration

- Comet **completely synchronizes** VMs on both sides (phone & server).
 - In Java, everything you need for program execution is stored in memory.
 - Program code, stack, heap, & CPU state
 - DSM can synchronize these.
- Any side can execute a thread, since they both know everything necessary for program execution.
 - The PC is synchronized, so both sides know the next instruction to execute.
 - The registers are synchronized, so they both know the CPU state.
 - The stack & the heap are synchronized, so they know the memory state.

Summary

- Memory sharing among threads
 - By default, they share the same address space
- Memory sharing among processes
 - Shared memory API & semaphore API
 - Virtual-physical memory mapping implements this.
- Memory sharing across machines
 - Write-update
 - Write-invalidate
- Memory sharing across threads on different machines
 - Use case: code offloading

Acknowledgements

- These slides contain material developed and copyrighted by Indranil Gupta (UIUC).